
Mixing Matters: Evidence Position Bias Across Attention and State-Space Language Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Language models often answer multi-document questions less accurately when
2 the relevant document is placed in the middle rather than near the beginning or
3 end. We test whether this pattern differs between models that use attention and
4 models from the Mamba family, which process context through a running state.
5 For each question, we move the answer-bearing document through ten positions
6 while keeping the question and nine distracting documents fixed. Our main com-
7 parisons use 800 exploratory questions; smaller task and prompt checks use 50 or
8 200 examples, and the larger held-out set has not been examined. Pythia-2.8B was
9 about five percentage points more accurate when the answer appeared near the be-
10 ginning than in the middle, whereas two similarly sized, Pile-trained Mamba mod-
11 els showed no such beginning advantage. A more closely matched 8B comparison
12 pointed in the same direction: a hybrid Mamba-2 model with a small number of at-
13 tention layers showed a larger beginning advantage than the pure Mamba-2 model,
14 but the difference between them remains too uncertain for a firm conclusion. The
15 contrast appears in larger but not the smallest model pairs, is not explained by the
16 tested change of training corpus, and reappears for Pythia on a synthetic retrieval
17 task. Prompt changes can strongly alter the pattern, and exploratory internal-state
18 analyses show that Mamba retains position information even without the begin-
19 ning advantage. These findings suggest that model design affects how available
20 evidence is used across a context, but they do not show that attention itself is the
21 cause.

22 1 Introduction

23 Transformers often read multi-document contexts unevenly: accuracy peaks when the answer-
24 bearing document is near an edge and falls in the middle (Liu et al., 2024). This behavior matters
25 because retrieval systems cannot assume that all evidence positions are equally usable. Long-context
26 benchmarks expose related failures across question answering, summarization, and retrieval (Bai et
27 al., 2024; Shaham et al., 2023; Li et al., 2024). Prior work links lost-in-the-middle behavior to
28 U-shaped attention patterns (Hsieh et al., 2024a), but model-family comparisons still confound se-
29 quence mixing with data, tokenizer, scale, depth, and positional encoding (Su et al., 2024; Press et
30 al., 2022). State-space models such as Mamba (Gu and Dao, 2024; Dao and Gu, 2024) provide a
31 useful contrast because they compress the prefix into a recurrent state instead of using dense causal
32 attention.

33 We make three contributions. First, we build an auditable ten-position evaluation with floor and
34 ceiling anchors, preserved generations, and inference settings recorded per example. Second, we
35 find a paired architecture-position interaction in a Pile-trained 2.8B comparison and test it with a
36 more tightly matched 8B pure/hybrid pair. Third, we organize scale, corpus, task, prompt, sink,

and probe results into mechanistic evidence with explicit uncertainty. All reported results use an exploratory split; the held-out confirmatory split remains unspent.

2 Related work

Position bias in long contexts Liu et al. (2024) established the lost-in-the-middle curve in retrieval and multi-document QA. This task derives its questions from Natural Questions (Kwiatkowski et al., 2019), while LongBench, ZeroSCROLLS, and LooGLE broaden long-context evaluation across tasks, languages, and reasoning demands (Bai et al., 2024; Shaham et al., 2023; Li et al., 2024). Subsequent work measures effective context length (Hsieh et al., 2024b), associates the curve with attention allocation (Hsieh et al., 2024a), and proposes reordering, compression, retrieval, and context-use interventions (An et al., 2024; Peysakhovich and Lerer, 2023; Jiang et al., 2024; Xu et al., 2024; Agrawal et al., 2024). Irrelevant context can itself distract a model, motivating our use of the same distractors at every tested position (Shi et al., 2023). We focus on whether the curve changes across sequence mixers.

Sequence mixers Dense self-attention (Vaswani et al., 2017) mixes token pairs under a causal mask and commonly encodes order with rotary embeddings or linear attention biases (Su et al., 2024; Press et al., 2022). Alternative long-sequence mixers include structured state spaces, recurrent weighted key-value models, and long convolutions (Gu et al., 2022; Peng et al., 2023; Poli et al., 2023). Mamba and Mamba-2 carry information through selective state-space recurrences (Gu and Dao, 2024; Dao and Gu, 2024). Hybrid models such as Griffin, Jamba, and Nemotron-H combine recurrent or state-space blocks with attention (De et al., 2024; Lieber et al., 2024; NVIDIA, 2025). Waleffe et al. (2024) release the pure and hybrid 8B Mamba-2 checkpoints used in our tightest controlled comparison.

Attention sinks and probes Xiao et al. (2024) showed that causal transformers can concentrate attention on initial tokens, and Sun et al. (2024) connected unusually large internal activations to a small set of tokens. We test whether this attention-sink statistic covaries with primacy, treating the result as correlational because a valid sink-blocking ablation is absent. Our position probe follows the principle that a diagnostic classifier needs a control task before its accuracy can be interpreted as evidence about a representation (Hewitt and Liang, 2019).

3 Measurement and statistics

Task and anchors We use the 2,655 released questions from the Lost-in-the-Middle 10-document QA setting (Liu et al., 2024), which is built from Natural Questions (Kwiatkowski et al., 2019). Each question has one gold document and nine distractors. We build ten prompts per question, moving only the gold document through positions 0 to 9 while retaining the same distractors and token budget. A no-gold floor measures guessing or memorized knowledge, and a gold-only ceiling measures answerability under perfect retrieval.

Edges and statistics Let a_S , a_M , and a_E be mean accuracy at start positions $\{0, 1\}$, middle positions $\{4, 5\}$, and end positions $\{8, 9\}$. We define primacy as $a_S - a_M$ and recency as $a_E - a_M$. The question is the unit of analysis. Each of 10,000 bootstrap samples draws complete ten-position question bundles with replacement, following the bootstrap principle of resampling the independent unit (Efron, 1979), and model comparisons reuse the same sampled questions. We report percentile 95% confidence intervals and apply Holm’s sequential correction to the primacy and recency tests within each comparison family (Holm, 1979). The fixed split contains 800 exploratory and 1,855 confirmatory questions. Only the exploratory split is reported.

Reproducibility and controls The harness vendors the official prompt builder and best_subspan_em scorer at a pinned commit. It also records normalized exact match, following the normalization convention popularized by SQuAD (Rajpurkar et al., 2016), and first-line extraction. Generation is greedy with temperature 0, at most 32 new tokens, fixed seeds, BF16 where supported, pinned model revisions, and one Mamba execution path. Each unique run preserves prompts, raw generations, failures, token counts, software versions, and decoding settings. The released key-value

task provides an executed positive control: the pipeline detects a +0.38 Pythia start-versus-middle effect. Negative and distractor-order controls are implemented but have no committed run artifacts, so we do not count them as passed controls.

4 Architecture comparisons

Primary 2.8B comparison We compare pythia-2.8b (Biderman et al., 2023) with mamba-2.8b and mamba2-2.7b. All are base models pretrained on the Pile (Gao et al., 2020), have similar parameter counts, and use related tokenizers. Table 1 and Figure 1 show a positive primacy edge for Pythia and no positive primacy edge for either Mamba model. The paired Mamba-minus-Pythia primacy interactions are -0.053 and -0.070 (both Holm $p < 10^{-4}$), while Mamba-1 versus Mamba-2 is uncertain ($+0.017$, 95% CI $[-0.005, 0.038]$). All three models have positive recency edges. These comparisons match corpus and approximate scale, but not depth, positional encoding, state size, or checkpoint family.

Table 1: Pile-trained models on 800 exploratory questions. Intervals use a question-level paired bootstrap; p is Holm-corrected.

Model	Primacy (95% CI)	Recency (95% CI)	Floor / Ceiling
pythia-2.8b	$+0.052$ $[0.032, 0.072]$, $p < 10^{-4}$	$+0.075$ $[0.054, 0.097]$, $p < 10^{-4}$	0.091/0.640
mamba-2.8b	-0.001 $[-0.016, 0.014]$, $p=0.91$	$+0.077$ $[0.058, 0.097]$, $p < 10^{-4}$	0.115/0.615
mamba2-2.7b	-0.018 $[-0.034, -0.003]$, $p=0.02$	$+0.105$ $[0.086, 0.124]$, $p < 10^{-4}$	0.128/0.619




Figure 1: Accuracy as the answer-bearing document moves through the context on 800 exploratory questions. Pythia rises at the beginning, while both smaller Mamba variants do not; the matched 8B hybrid also rises more than the pure model, but their paired difference remains uncertain. Dotted lines show each model’s no-gold floor and gold-only ceiling.

Matched 8B attention comparison Our strongest control uses NVIDIA’s pure and hybrid 8B Mamba-2 checkpoints (Waleffe et al., 2024), which share training data, tokenizer, scale, depth, and positional setup. The hybrid contains 7% attention layers. Within models, the pure model’s primacy edge is uncertain ($+0.008$, $p = 0.36$) and the hybrid’s is positive ($+0.027$, $p = 0.007$). That difference of significance is not itself a significant difference. The paired hybrid-minus-pure primacy effect is $+0.019$ (95% CI $[-0.006, 0.044]$, Holm $p = 0.144$), while the recency effect is -0.028 (95% CI $[-0.055, -0.001]$, Holm $p = 0.089$). This direction supports the 2.8B observation but does not isolate a causal attention effect.

Table 2: Matched 8B primacy comparison on 800 exploratory questions. The paired interaction is the relevant architecture test.

Contrast	Primacy	95% CI	Holm p
pure Mamba-2	+0.008	$[-0.009, 0.024]$	0.362
hybrid	+0.027	$[0.008, 0.047]$	0.007
hybrid minus pure	+0.019	$[-0.006, 0.044]$	0.144

106 5 Scope checks

107 **Scale** Across five approximate size pairs, the Pythia-minus-Mamba primacy gap is near zero at
 108 the two smallest scales and is +0.062, +0.069, and +0.053 in the three larger pairs (Figure 2).
 109 The smallest Pythia ceiling is only 0.009, so this is a descriptive capability-and-family trend, not
 110 evidence for a depth threshold.

111 **Training corpus** Holding the Mamba-2.8B architecture fixed, the Pile-minus-SlimPajama primacy
 112 effect is -0.007 (95% CI $[-0.030, 0.016]$, $p = 0.568$) and the recency effect is -0.016 (95% CI
 113 $[-0.044, 0.012]$, Holm $p = 0.502$) (Soboleva et al., 2023). The floor differs substantially (0.270
 114 versus 0.114), so this comparison finds no edge-shape change but does not rule out training-data
 115 effects (Figure 2).

116 **Synthetic retrieval** On 50 RULER niah_single_1 examples at 2K tokens (Hsieh et al., 2024b),
 117 Pythia again has a primacy edge (+0.12, 95% CI $[0.05, 0.20]$). Both Mamba models score 1.0 at
 118 every depth, so this check reproduces the transformer pattern but cannot compare mixer curves on
 119 the synthetic task (Figure 3).

120 **Descriptive system comparison** We also run Llama-3.1-8B (Dubey et al., 2024), Qwen2.5-7B
 121 (Yang et al., 2024), and Nemotron-H-8B (NVIDIA, 2025). All three show positive primacy edges
 122 (+0.076, +0.041, +0.067). Only Llama versus Qwen has a significant paired primacy interaction
 123 (+0.036, Holm $p = 0.016$). Because the systems differ on architecture, data, tokenizer, token count,
 124 depth, and positional encoding, these results establish prevalence, not an architectural explanation
 125 (Figure 4).

126 6 Mechanistic evidence and intervention

127 **Attention sinks track primacy** Across the Pythia family, the 160M model has almost no token-
 128 0 attention sink (final-layer share 0.003), while the 2.8B model reaches 0.455 (Xiao et al., 2024).
 129 Pythia-410M has a small final-layer sink (0.034) and no significant primacy edge. Figure 3 is con-
 130 sistent with a sink mechanism, but the measurements share model scale as a confound and no valid
 131 sink ablation was completed.

132 **Utilization, not storage, is a hypothesis** A balanced linear probe on frozen mid-depth states pre-
 133 dicted edge versus middle position above its shuffled-label control for Pythia (0.603 versus 0.484)
 134 and Mamba (0.648 versus 0.508). Thus position information remains linearly decodable even when
 135 Mamba has no positive primacy edge. This supports a utilization hypothesis, but probe decodability
 136 does not prove that the model stores or uses the answer-bearing content.

137 **Training-free prompt intervention** On 200 exploratory questions, repeating the question both be-
 138 fore and after the documents increases Mamba primacy from 0.000 to +0.103 ($p < 10^{-4}$). Placing
 139 the question only before the documents does not create primacy (+0.020, $p = 0.366$), but reduces re-
 140 cency from +0.070 to +0.010. The intervention is training-free and practically useful, but it changes
 141 query repetition and token placement together, so it is not a clean test of fixed-state compression.

142 **Prompt sensitivity** Changing the prompt template reduces Pythia primacy from +0.070 to +0.005
 143 in the 200-question mechanism subset. This result warns against treating the measured magnitude
 144 as a template-invariant architectural constant.

7 Discussion and limitations

The most secure result is the 2.8B exploratory architecture-position interaction. The matched 8B direction, sink correlation, and probe results jointly make attention-mediated utilization a plausible hypothesis, while prompt sensitivity shows that architecture is not the only determinant.

Several limitations set the confirmatory boundary. All analyses use the exploratory split, so the planned 1,855-question test must be run under a frozen analysis before confirmatory language is warranted. The 2.8B and scale comparisons retain depth, positional-encoding, state-size, and capability confounds. The matched 8B paired contrast crosses zero after Holm correction, and its run manifest records a dirty source tree, weakening exact reproducibility. The negative and distractor-order controls lack committed outputs. A 50-generation blinded audit sample exists, but human labels were not completed. Finally, the attempted hybrid sink block did not alter the model because its custom attention path did not honor the hook; it is not an ablation result. The next decisive study should freeze exclusions and contrasts, complete these controls, repair the hybrid intervention, and evaluate the held-out split once.

8 Conclusion

On 800 exploratory Lost-in-the-Middle questions, Pythia-2.8B and two Pile-trained Mamba models have different primacy profiles under a paired ten-position evaluation. A matched 8B comparison is directionally consistent but statistically uncertain, so the present evidence supports an architecture-position interaction without establishing that attention causes it. The held-out split and the remaining control work provide a clear path from this exploratory workshop result to a confirmatory study.

A Experiment artifacts

The repository preserves raw generations, environment records, machine-readable summaries, and plots for every reported phase. The panels below reproduce the exact committed artifact images without redrawing or smoothing them. They summarize scope and mechanism checks that support interpretation but do not upgrade the exploratory results to confirmatory or causal evidence.




Figure 2: Exploratory scope checks. The scale sweep is descriptive because model capability and family differences vary together. The corpus comparison holds the Mamba architecture fixed and finds no clear change in curve shape, despite a substantial change in overall accuracy.




Figure 3: Exploratory task and mechanism artifacts. Pythia has a middle dip on the 2K synthetic task, while both Mamba models saturate and therefore cannot reveal a comparable curve. Late-layer token-0 attention grows with Pythia scale and covaries with the beginning advantage, but scale is a shared confound and no valid causal ablation was completed.




Figure 4: Phase 8 descriptive production-system comparison. All three systems show a beginning advantage, but differences in architecture, data, tokenizer, training, and capability prevent attribution to a single design choice.

References

- D. Agrawal, S. Gao, and M. Gajek. Can't remember details in long documents? You need some R&R. In *Findings of the Association for Computational Linguistics: EMNLP*, pages 12692-12704, 2024.
- C. An, S. Huang, J. Zhang, et al. Make your LLM fully utilize the context. In *Advances in Neural Information Processing Systems*, 2024.
- Y. Bai, X. Lv, J. Zhang, et al. LongBench: A bilingual, multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 3119-3137, 2024.
- S. Biderman, H. Schoelkopf, Q. Anthony, et al. Pythia: A suite for analyzing large language models across training and scaling. In *International Conference on Machine Learning*, pages 2397-2430, 2023.
- T. Dao and A. Gu. Transformers are SSMs: Generalized models and efficient algorithms through structured state space duality. In *International Conference on Machine Learning*, 2024.
- S. De, S. L. Smith, A. Fernando, et al. Griffin: Mixing gated linear recurrences with local attention for efficient language models. *arXiv:2402.19427*, 2024.
- A. Dubey, A. Jauhri, A. Pandey, et al. The Llama 3 herd of models. *arXiv:2407.21783*, 2024.
- L. Gao, S. Biderman, S. Black, et al. The Pile: An 800GB dataset of diverse text for language modeling. *arXiv:2101.00027*, 2020.
- B. Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1-26, 1979.
- A. Gu and T. Dao. Mamba: Linear-time sequence modeling with selective state spaces. In *Conference on Language Modeling*, 2024.
- A. Gu, K. Goel, and C. Ré. Efficiently modeling long sequences with structured state spaces. In *International Conference on Learning Representations*, 2022.
- J. Hewitt and P. Liang. Designing and interpreting probes with control tasks. In *Proceedings of EMNLP-IJCNLP*, pages 2733-2743, 2019.
- S. Holm. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65-70, 1979.
- C.-P. Hsieh, S. Sun, S. Krizan, S. Acharya, D. Rekes, F. Jia, and B. Ginsburg. Found in the middle: Calibrating positional attention bias improves long context utilization. In *Findings of the Association for Computational Linguistics: ACL*, pages 14982-14995, 2024.
- C.-P. Hsieh, S. Sun, S. Krizan, et al. RULER: What's the real context size of your long-context language models? In *Conference on Language Modeling*, 2024.
- H. Jiang, Q. Wu, X. Luo, et al. LongLLMLingua: Accelerating and enhancing LLMs in long context scenarios via prompt compression. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 1658-1677, 2024.
- T. Kwiatkowski, J. Palomaki, O. Redfield, et al. Natural Questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:452-466, 2019.
- O. Lieber, B. Lenz, H. Bata, et al. Jamba: A hybrid Transformer-Mamba language model. *arXiv:2403.19887*, 2024.
- J. Li, M. Wang, Z. Zheng, and M. Zhang. LooGLE: Can long-context language models understand long contexts? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 16304-16333, 2024.
- N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157-173, 2024.
- NVIDIA. Nemotron-H: A family of accurate and efficient hybrid Mamba-Transformer models. *arXiv:2504.03624*, 2025.
- A. Peysakhovich and A. Lerer. Attention sorting combats recency bias in long context language models. *arXiv:2310.01427*, 2023.

218 B. Peng, E. Alcaide, Q. Anthony, et al. RWKV: Reinventing RNNs for the Transformer era. In *Findings of the*
219 *Association for Computational Linguistics: EMNLP*, pages 14048-14077, 2023.

220 M. Poli, S. Massaroli, E. Nguyen, et al. Hyena hierarchy: Towards larger convolutional language models. In
221 *International Conference on Machine Learning*, volume 202 of PMLR, pages 28043-28078, 2023.

222 O. Press, N. A. Smith, and M. Lewis. Train short, test long: Attention with linear biases enables input length
223 extrapolation. In *International Conference on Learning Representations*, 2022.

224 P. Rajpurkar, J. Zhang, K. Lopyrev, and P. Liang. SQuAD: 100,000+ questions for machine comprehension of
225 text. In *Proceedings of EMNLP*, pages 2383-2392, 2016.

226 U. Shaham, M. Ivgi, A. Efrat, J. Berant, and O. Levy. ZeroSCROLLS: A zero-shot benchmark for long text
227 understanding. In *Findings of the Association for Computational Linguistics: EMNLP*, pages 7977-7989,
228 2023.

229 F. Shi, X. Chen, K. Misra, N. Scales, D. Dohan, E. H. Chi, N. Schärli, and D. Zhou. Large language models can
230 be easily distracted by irrelevant context. In *International Conference on Machine Learning*, volume 202 of
231 PMLR, pages 31210-31227, 2023.

232 D. Soboleva, F. Al-Khateeb, R. Myers, et al. SlimPajama: A 627B token cleaned and deduplicated version of
233 RedPajama. Cerebras technical report, 2023.

234 J. Su, Y. Lu, S. Pan, A. Murtadha, B. Wen, and Y. Liu. RoFormer: Enhanced transformer with rotary position
235 embedding. *Neurocomputing*, 568:127063, 2024.

236 M. Sun, X. Chen, J. Kolter, and Z. Liu. Massive activations in large language models. *arXiv:2402.17762*, 2024.

237 A. Vaswani, N. Shazeer, N. Parmar, et al. Attention is all you need. In *Advances in Neural Information*
238 *Processing Systems*, 2017.

239 R. Waleffe, W. Byeon, D. Riach, et al. An empirical study of Mamba-based language models.
240 *arXiv:2406.07887*, 2024.

241 G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis. Efficient streaming language models with attention sinks. In
242 *International Conference on Learning Representations*, 2024.

243 A. Yang, B. Yang, B. Zhang, et al. Qwen2.5 technical report. *arXiv:2412.15115*, 2024.

244 F. Xu, W. Shi, and E. Choi. RECOMP: Improving retrieval-augmented LMs with compression and selective
245 augmentation. In *International Conference on Learning Representations*, 2024.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 identify the exploratory architecture interaction, the uncertain matched 8B contrast, and the correlational status of the mechanism evidence.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 states the exploratory-only scope, the residual confounds, the uncertain paired 8B result, the dirty producing tree, the missing control outputs and audit labels, and the invalid sink intervention.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical and contains no theoretical results.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [No]

Justification: Section 3 discloses the dataset, split, inference settings, scorer, and analysis, but the matched 8B run manifest records uncommitted producing-code changes, so exact reconstruction of that result is not yet guaranteed.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The dataset and checkpoints are public and the repository contains the harness, runbooks, and raw outputs, but a fully anonymized submission archive has not yet been prepared.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 3 gives the split, the fixed generation settings, and the edge and bootstrap definitions; no training is performed beyond the frozen-state linear probe, whose layer is fixed a priori.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Every edge and contrast reports a 95 percent question-level paired-bootstrap interval (10,000 resamples) with Holm-corrected p -values, defined in Section 3.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The artifacts identify single-GPU hardware including L40S and T4 runs and preserve per-run configurations, but the paper does not yet report elapsed time and total project compute for every experiment.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

Answer: [Yes]

Justification: The work is an analysis of released public models on a released public benchmark, with no human subjects and no new data collection.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The work is a diagnostic analysis of position bias in existing models; it introduces no deployable artifact with a direct path to societal harm.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases analysis code and evaluation outputs, not any high-risk model or scraped dataset.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The dataset, scorer, and every model checkpoint are cited to their original publications; all are used under their public release terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The released harness and per-phase runbooks document the sweep, scoring, and analysis commands that produce every reported number.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The paper involves no crowdsourcing and no human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The paper involves no human subjects and therefore required no IRB review.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: Large language models are the objects of study, not a component of the method; no LLM was used as an original or non-standard part of the methodology.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.